



NHID-Clinical

FRAMEWORK V1.3 · SOURCE COMMIT E6EA815

NHID-Clinical Playbook

Evaluation, Implementation, and Evidence Guide

Version 1.0 · June 2026 · CC BY 4.0

Author: Brianna Baynard · Independent, practitioner-led

NIST Public Comment: NIST-2025-0035-0026

OPEN PROPOSAL — NOT A STANDARD OR REGULATORY REQUIREMENT

NHID-Clinical is an early-stage open proposal by Brianna Baynard. It is a voluntary open proposal — not an accredited standard, not a certification program, and not a regulatory requirement. No HIPAA compliance, security guarantees, or liability protection are implied.

Evaluation, Implementation, and Evidence Guide

Playbook version	1.0
Published	2026-09-04 · revised 2026-09-05
Framework version	NHID-Clinical v1.3
Source commit	e68a65d — the commit the evidence below was measured and clean-clone verified at
Canonical location	docs/NHID-Clinical-Playbook.md in NHID-Clinical/NHID-Clinical
Licence	CC BY 4.0
Author	Brianna Baynard — independent, practitioner-led
Provenance	Every substantive section traces to a source listed in playbook-source-inventory.md

What this document is not. NHID-Clinical is a voluntary open proposal. It is **not** an accredited standard, a certification programme, or a regulatory requirement. Publishing this Playbook does not make it any of those things. No HIPAA compliance, security guarantee, or liability protection is implied or conferred. Regulatory mappings in Part V are *mappings*, not legal advice and not compliance claims. **There are no known production deployments, customers, design partners, or external validations of NHID-Clinical.** Where this document describes what an organisation would do, it describes a method, not an observed practice.

How to use this Playbook

The website answers *what is this and should I care*. The repository answers *show me the code*. ****This Playbook answers *how exactly do I do this*****

If you are...	Start at
Deciding whether this is relevant at all	Part I
Evaluating an existing AI voice workflow	Part I → Part III
Building or integrating an implementation	Part II → Part IV
Assessing evidence, or preparing for review	Part I → Part V
Looking for a schema, checklist, or example	Appendices

Part I is self-contained. A reader who stops after it should understand the framework, its scope, its evidence, and its limits without reading further.

A convention used throughout. Every capability is labelled:

Label	Meaning
Implemented	Exists in the repository, exercised by the test suite
Reference implementation	Exists and works, offered as a reference rather than a product
Conceptual	Specified or designed; not built
Future	Named as a direction; neither specified nor built
Unknown	The repository does not establish this. Requires human judgment

Nothing conceptual or future is described as though it currently works.

Part I — Executive Brief

1. What NHID-Clinical is

NHID-Clinical is a **voluntary behavioural baseline** for AI voice agents on B2B healthcare administrative calls, together with a conformance test suite and a deterministic reference policy engine.

It defines five controls covering what an AI agent should do on a call — disclose that it is automated, withhold protected data until it has, avoid passing as human, honour escalation to a person, and leave an audit record — and it makes those controls **testable**: the same transcript produces the same pass/fail result every time, with a reason code.

It is a *proposal*. Its authority comes from being checkable, not from adoption or endorsement.

2. The operational problem

On a payer-provider administrative call, operational data starts moving early. A caller asks for a member ID, an NPI, a date of birth, a claim number — often within the first few turns.

Telephony authenticates the *number*. Identity systems authenticate the *account*. Neither establishes that **this caller is an automated system**, nor that it is allowed to act for the organisation it names.

The gap between an AI agent beginning to operate and the receiving party being able to know it is non-human is what this framework calls **impersonation latency**. Everything exchanged inside that window is exchanged with a counterparty whose nature the receiver could not assess.

NHID-Clinical addresses that specific window. It does not address fairness, clinical safety, or model quality — those are deliberately out of scope.

3. Who it is for

Audience	Use
Payer operations and compliance	Evaluate what AI callers currently establish, before changing anything (Part III)
Provider organisations and their vendors	Understand what a receiving party may reasonably expect (Parts I–II)
AI voice vendors	Self-assess against a published, deterministic suite (Parts II, IV)
Procurement and vendor risk	RFP language, a trust questionnaire, and an evidence model (Part V, appendices)
Governance, risk, audit	Evidence taxonomy, audit event model, regulatory mapping (Part V)

4. Scope

In scope: B2B administrative voice workflows — AI systems calling payer offices on behalf of providers, vendors, or plan administrators; and payer-side agents calling providers. Eligibility, prior authorisation, claim status, benefits verification.

Observable behaviour on the call, and the evidence record it leaves.

5. Non-scope

Explicitly **not** covered, and not claimed:

- Patient-facing calls, clinical decision support, internal tooling
- Fairness, bias, clinical safety, model quality — see [scope-boundary-fairness-clinical.md](#)

- Caller identity *verification* (the framework standardises disclosure and trace behaviour; it does not verify who the caller is, and does not certify vendors)
- Legal compliance of any kind
- Voice biometrics, call recording policy, consent capture

6. Relationship to existing infrastructure

NHID-Clinical sits **above** carrier authentication and **beside** identity management. It replaces neither.

Layer	What it establishes	What it does not
STIR/SHAKEN	The calling number is legitimate	Whether the caller is automated, or authorised for PHI
Identity / IAM	An account is authenticated	What that account does on a voice call
NHID-Clinical v1.3	The agent disclosed it is automated before PHI moved, and left a record	Who the caller is
NHID-Auth v2 (<i>reference implementation</i>)	A cryptographically signed, scoped delegation from a provider	—

7. The five canonical controls

Five controls. Four are deterministic behavioural controls governing what the agent does on the call; the fifth is the audit and evidence control, which changes nothing about behaviour and records what the other four did.

Control	Name	Kind	In one line
IDG-01	Identity Disclosure Gate	Behavioural	The agent identifies itself as automated before any PHI exchange
PDX-01	Pre-Data Exchange Gate	Behavioural	No PHI — member ID, NPI, DOB, claim number — until IDG-01 disclosure is confirmed
DBC-01	Deceptive Behaviour Check	Behavioural	No claim of human identity; no artifacts designed to imply human presence
EIT-01	Escalation Implementation Test	Behavioural	A human escalation path is communicated and honoured on request
ATR-01	Audit Trail	Audit / evidence	A machine-readable event trace: disclosure timestamps, state transitions, escalation events, execution context

DLG-01 (Delegated Authority) is opt-in and is not one of the five. Absent a delegation it returns `DLG01_NOT_EVALUATED` and changes nothing.

Part II gives each control's inputs, expected behaviour, failure behaviour, evaluation logic, and limitations.

8. Current evidence and validation status

Four separate bodies of evidence. **They measure different things and are not interchangeable.** Collapsing them into one "accuracy" number would misrepresent all four.

Evidence body	What it measures	Result	Date
Conformance suite	Technical test execution against the engine, adapters, API and invariants	1094 collected · 1094 executed · 1094 passed · 0 failed, skipped, xfailed, xpassed	2026-09-04, commit e68a65d
Fabricate Battle-Test Corpus	Detection against 550 real-world voice AI conversations, 127 of them compliant	IDG-01 70/70 · PDX-01 41/41 · DBC-01 183/200 (91.5%) · EIT-01 169/171 (98.8%). False positives on clean conversations: 0, 0, 5, 5 of 127	CI-gated, unchanged
Governance Evaluation Corpus	Detection of labelled governance conditions across 25 scenarios / 55 turns	30 of 32 = 93.8% . False positives 0 of 5 compliant scenarios. 12 unexpected detections on violation scenarios, reported separately	2026-09-04
Adversarial corpus	Robustness against 40 deliberately hostile scenarios	See safety/adversarial-testing-report.md	—

These are four different denominators. 1094/1094 is a *test pass rate*, not a detection rate. 93.8% is a *detection rate*, not a test pass rate. Neither is an accuracy figure for the framework as a whole.

On the 12 unexpected detections. The published false-positive figure is measured only over the five compliant scenarios, so a control firing where it was not expected on any of the other twenty could not appear in it. There are twelve such detections — eight found when the quantity was first measured, and four more added by the G2 decision, which made a bare organisational name an IDG-01 violation on four scenarios that do not declare one. They are reported as a distinct quantity rather than folded into the false-positive rate, because a compliant scenario emitting anything is an engine defect, whereas a violation scenario emitting an undeclared rule is usually the corpus under-specifying its own turns. Judging which requires reading the scenario. See [governance-corpus-remediation.md](#).

9. Current limitations and unresolved questions

Stated plainly, because a reader who discovers these later has been misled.

What is not established, at all:

	Status
Production deployments	None known
Customers, design partners, pilot participants	None
Independent external validation or audit	None
Certification or accreditation	None, and none sought
Behaviour under production load	Unknown — not load-tested
Regulatory endorsement	None

Known technical limitations:

- **Detection is lexical.** PDX-01 and DBC-01 match phrase maps plus a small number of structural rules. A paraphrase outside the map is missed.
- **DBC-01 is the least precise control** — 91.5% detection with 5 false positives on 127 clean conversations.
- ****Disclosure *sufficiency* is judged only on the disclosing turn.**** Since the G2 decision, IDG-01 requires the disclosing turn to state a non-human identity affirmatively, so a bare organisational name *is* flagged. The check needs a harness that sets [disclosure_established_prior](#); without it the permissive default applies and only contradiction is caught.

- **Sequencing checks need a cooperating harness.** Same-turn disclosure-and-request detection requires `disclosure_established_prior`; absent it the check does not run.
- **ATR-01 persistence is external.** The engine emits the audit trail; it cannot detect that a downstream store failed to persist it.
- **Escalation quality is not assessed.** EIT-01 verifies that an escalation path exists and was honoured, not that a competent human answered.

Four open governance questions (G1–G4) are unresolved and are *not* silently decided anywhere in this Playbook. They are stated in Part V §5 and referenced wherever they bear on a control.

Part II — Framework

This part assembles canonical material rather than restating it. The normative sources are [CONTROL_DECISION_TABLE.md](#), [enforcement-profile.md](#) and [NHID_AUDIT_EVENT_SPEC_v1.0.md](#). Where this Playbook and a normative source disagree, the normative source wins and this document is wrong.

1. Terminology

Term	Meaning
Non-human actor	Any automated participant in a call. Preferred over "bot"
Interaction boundary	The point at which an automated participant begins operating with a counterparty
Impersonation latency	The interval between an agent beginning to operate and the counterparty being able to establish that it is non-human
Disclosure	An assertion, by the agent, that it is automated
Turn	One utterance-and-response unit of a call; the granularity at which controls evaluate
PHI	Protected health information. In practice: member ID, NPI, date of birth, claim number
Control	One of the five canonical governance checks
Violation	A BoundaryViolation — rule_id , description , severity
PolicyDecision	The single output of evaluating all controls against one turn
CAS	Call Authorization Score — a downstream routing signal, not an evaluator
Tier 0	Observe-only shadow evaluation. Decisions recorded, never enforced

Full list: [docs/terminology.md](#).

2. The two-layer model

NHID-Clinical separates **the governance layer** (the five controls, which evaluate behaviour) from **the authorization layer** (NHID-Auth v2, which verifies cryptographic delegation). The governance layer works without the authorization layer. The reverse is not true.

3. The control model

Evaluation is **pure and deterministic**. [evaluate_all\(session, event, delegation=None\)](#) returns exactly one [PolicyDecision](#), performs no I/O, and never raises. The same inputs always produce the same decision and the same reason code.

3.1 [PolicyDecision](#) — the output contract

Field	Meaning
action	The single PolicyAction the receiver MUST execute for this turn
reason_code	Stable machine token identifying which condition produced the action
violations[]	Every BoundaryViolation raised by all controls this turn — merged, never collapsed
next_state	Advisory workflow state label
policy_version	Engine version that produced the decision

A [PolicyDecision](#) carries no score. CAS is computed separately, downstream.

3.2 [PolicyAction](#) — exactly five

Action	Receiver obligation	Blocks PHI?
DISCLOSE_IDENTITY	MUST require identity disclosure before material interaction continues	Yes
DENY_DATA	MUST NOT request, accept, or release PHI on this turn	Yes
ESCALATE_HUMAN	MUST provide a functional human escalation path	Indirect
LOG_ONLY	MUST record the finding; SHOULD route to human review by severity	No
CONTINUE_AI	None beyond proceeding	No

A sixth action must not be introduced — a new action would be a sixth control in disguise.

3.3 The enforcement ladder

A single turn may trip several controls, but a receiver needs one unambiguous action. When outcomes conflict, the most-protective action wins by fixed precedence:

DENY_DATA > ESCALATE_HUMAN > DISCLOSE_IDENTITY > LOG_ONLY > CONTINUE_AI

The ladder **selects an action; it does not suppress findings**. Every control still evaluates independently and contributes its violations. The ladder never re-decides conformance.

3.4 The CAS authority boundary

The Call Authorization Score is a downstream assessment and routing mechanism, derived after and from the [PolicyDecision](#).

CAS may trigger human review and influence queue priority. **CAS must not** override a [PolicyDecision](#), convert a restrictive action into allowed access, or independently determine conformance.

Normative invariant. The Enforcement Profile SHALL consume [PolicyDecision](#) outputs and SHALL NOT independently evaluate control conformance. CAS SHALL be derived from, and downstream of, the [PolicyDecision](#). The five controls remain the sole source of conformance decisions.

CAS is a **research component**. It must not be presented as a compliance score.

4. The five controls in detail

The authoritative per-control reference is [CONTROL_DECISION_TABLE.md](#), which gives trigger, pass condition, fail condition, evidence location, limitation, test coverage and corpus status for each. Summarised here with the parts an implementer most often needs.

IDG-01 — Identity Disclosure Gate

Intent	The agent discloses non-human identity before sensitive interaction
Trigger	Any turn where <code>event_type == IDENTITY_DISCLOSURE</code> or <code>contains_phi == True</code>
Pass	<code>disclosure_timestamp</code> set, <code>identity_assertion_text</code> non-empty, and on the disclosing turn the assertion neither introduces a human persona without stating non-human identity nor denies being automated
Fail	<code>IDG01_DISCLOSURE_MISSING</code> · <code>IDG01_ASSERTION_TEXT_MISSING</code> · <code>IDG01_DISCLOSURE_CONTRADICTED</code>
Action	DISCLOSE_IDENTITY
Limitation	Judges disclosure <i>content</i> only for contradiction, not adequacy. Content checks apply only to the disclosing turn and require a harness that sets <code>disclosure_established_prior</code> . Does not verify the counterparty understood
Evidence	Fabricate 70/70, 0 FP of 127. Governance corpus 6/8 (75.0%)

Open question G2 bears directly on this control: a bare organisational name ("claims system") is not flagged, because the evaluation corpus labels that shape a violation in one scenario and does not in a structurally identical one. No rule separates them, so the engine abstains. See Part V §5.

PDX-01 — Pre-Data Exchange Gate

Intent	Prevent PHI exchange before valid disclosure
Trigger	Any turn where <code>contains_phi == True</code>
Pass	No protected-data exchange, or disclosure established on a prior turn
Fail	<code>PDX01_PHI_GATE_TRIGGERED</code> (no disclosure at all) · <code>PDX01_DISCLOSURE_NOT_SEQUENCED</code> (same turn as initial disclosure) · <code>PDX01_SCOPE_NOT_AUTHORIZED</code> (outside a verified DLG-01 scope)
Action	<code>DENY_DATA</code>
Limitation	Detection is lexical plus one structural rule; paraphrase outside the phrase map is missed. Sequencing needs <code>disclosure_established_prior</code>
Evidence	Fabricate 41/41, 0 FP of 127. Governance corpus 6/6 (100%)

Open question G4: whether same-turn disclosure followed by a PHI request should be a violation at all. The engine says yes; four corpus scenarios disagree by omission.

DBC-01 — Deceptive Behaviour Check

Intent	Prevent concealment of AI identity or impersonation of a human
Pass	<code>deceptive_artifact_flags == []</code>
Fail	Tier A — any explicit artifact flag (conceal request, impersonation, misrepresentation, evasive, ambiguous identity language, human-passing attempt). Tier B — impersonation phrase in the identity assertion, including first-person licensed-clinical role claims. Tier C — corpus-mined implied-humanity cues; weak disfluency cues require two or more
Action	<code>LOG_ONLY</code> — a deception finding is recorded and routed, not itself a PHI gate . Blocking on a deceptive turn comes from a co-occurring IDG-01/PDX-01 failure
Limitation	Tier C is inferential and stays active even when the same assertion discloses. Suppressing it after disclosure was tried and cost four real Fabricate detections — disclosing once and then passing as staff is a pattern the corpus labels deceptive. Third-person clinical references are deliberately not matched
Evidence	Fabricate 183/200 (91.5%), 5 FP of 127 — the least precise control . Governance corpus 9/9 (100%)

EIT-01 — Escalation Implementation Test

Intent	Escalation requests are honoured within a bounded window
Trigger	Any turn where <code>escalation_requested == True</code>
Pass	No escalation requested, or a path is available, or fulfilment is recorded — an <code>escalation_timestamp</code> together with an honouring <code>escalation_outcome</code>
Fail	Outcome <code>DEFLECTED</code> or <code>IGNORED</code> , or requested with no outcome within 5 turns
Action	<code>ESCALATE_HUMAN</code>

Limitation	Does not validate escalation <i>quality</i> — whether a real, competent human answered. Requests are keyword-detected, so the agent's own confirmation line matches as a request; recorded fulfilment is what prevents that reading as failure. An honouring outcome with no timestamp is a claim, not a record , and does not clear the control
Evidence	Fabricate 169/171, 5 FP of 127. Governance corpus 8/8 (100%), 0 false positives

ATR-01 — Audit Trail (the audit and evidence control)

Intent	Policy-relevant events are persisted for audit
Trigger	Every turn processed
Pass	Required audit fields present and non-empty, including execution context
Fail	<code>ATR01_AUDIT_FIELDS_MISSING</code>
Action	<code>LOG_ONLY</code>
Limitation	Persistence is outside the engine. The engine emits the trail; it cannot detect that an external store failed to write it
Evidence	Operational across evaluated sessions; persistence tested separately

Open question G3: whether ATR-01 should be evaluated from transcripts at all. It validates an audit *record*, and a transcript is not one.

5. The audit model

ATR-01's output is governed by `NHID_AUDIT_EVENT_SPEC_v1.0.md`, which defines the event schema, four event classes (`GOVERNANCE_DECISION`, `RULE_VIOLATION`, `DATA_ACCESS`, `ESCALATION_ACTION`), append-only guarantees, tamper-evidence via hash chaining, and retention.

Emission to **HL7 FHIR R4 AuditEvent** is **implemented** and validated in CI against the official validator. Mapping detail: `fhir-auditevent-mapping.md`.

6. Delegated authority — NHID-Auth v2

Status: reference implementation, opt-in, not part of the five controls.

NHID-Auth v2 addresses a question the behavioural controls cannot: *is this agent authorised to act for the provider it names?* It uses Ed25519-signed, scoped, expiring delegations verified offline.

DLG-01 evaluates a delegation when one is supplied. **Absent a delegation it returns `DLG01_NOT_EVALUATED` and changes nothing** — every pre-existing caller keeps its behaviour exactly.

Detail: `nhid-clinical-technical-specification.md` §7–9, `nhid-auth-pki-and-oauth2-integration.md`, and `framework/nhid-auth.html`.

Part III — Shadow Evaluation

Shadow evaluation is observe-only. You capture what your AI callers already do, replay it through the policy engine, and read the result. **The production call flow is not changed, not intercepted, and not gated.** Nothing an evaluation produces reaches a caller.

There is no prescribed duration. What follows is an **ordered sequence**, not a calendar. Each stage depends on the one before it; none has a required length, and an evaluation may legitimately stop after any of them. Any 30/60/90-day framing you encounter elsewhere in this repository is superseded by this statement.

1. Purpose — and what it can and cannot establish

Shadow evaluation can establish:

- Whether AI callers disclose, and *when* relative to the first data request
- Whether protected data moved before disclosure, and which fields
- Whether escalation requests were honoured
- Which controls fail, on which calls, with a reason code
- A baseline you can re-measure against later

Shadow evaluation cannot establish:

- That a vendor is compliant with any law or regulation
- That an agent is safe, accurate, or clinically appropriate
- That a human on the other end understood the disclosure
- **ATR-01 conformance.** Audit-trail completeness cannot be exercised from transcript replay — a replay harness constructs complete audit envelopes by definition. Audit completeness is assessed against a real event pipeline, not a transcript. This is the same limitation recorded as open question **G3**
- **DBC-01 Tier A (voice artifacts)**, unless your stack produces voice-forensics flags. Text heuristics still run

Stating these first is deliberate. An evaluation whose limits are discovered afterwards produces findings nobody trusts.

2. Prerequisites

Data	Call transcripts or turn-level logs from existing AI voice traffic. No audio required
Fields	Per turn: speaker, text, timestamp, and whichever protected fields were requested. Full schema: pilot-kit/minimal-event-schema.json
Software	The reference implementation and pilot-kit/measure_pilot.py . No account, no hosted service
Access	Read access to call logs. No production change of any kind

Sanity-check the toolchain before touching real data:

```
pip install -r requirements.txt
python docs/pilot-kit/measure_pilot.py --demo
```

3. Governance, approvals and data handling

Shadow evaluation reads call logs that may contain PHI. Treat it as any other secondary use of that data.

Concern	Consideration
Roles	Name an owner for the evaluation, a data owner for the logs, and a reviewer for findings. One person may hold more than one role

Concern	Consideration
Approval	Whatever your organisation requires for secondary use of call data. This framework asserts no view on what that is
Minimisation	The engine needs turn text and which protected fields were requested — not the field <i>values</i> . Redact values at capture where you can
Location	The reference implementation runs locally. Nothing is transmitted to the project
Retention	Set a retention period for captured turns before you capture them, not after
Access	Findings name failing calls. Scope access accordingly

This is not legal or privacy advice. Whether an evaluation is permissible under your obligations is a determination for qualified people in your organisation. Marked **unknown** by this framework, deliberately.

4. The evaluation sequence

Stage 1 — Workflow inventory

Establish what you actually have before measuring it.

- Which inbound or outbound call workflows involve AI agents?
- Which vendors or internal systems place them?
- Which involve protected data, and which fields?
- Roughly what volume, and over what period are logs retained?

Output: a list of candidate workflows, each with a data-sensitivity note. **Stop condition:** you can name every workflow in scope.

Stage 2 — Interaction capture

Extract turn-level records into the capture schema. Flat by design, so it can be filled from ordinary call logs.

Output: captured turns, one file per workflow. **Stop condition:** a sample replays through `--demo` without schema errors.

Stage 3 — Evidence collection

Where you have a real event pipeline (not just transcripts), collect the audit events alongside. **This is the only path to any ATR-01 signal** — see §1.

Output: event records, where they exist. **Stop condition:** you know whether you have them. "We do not" is a valid answer.

Stage 4 — Control evaluation

Replay captured turns through the real policy engine.

```
python docs/pilot-kit/measure_pilot.py --input captured_turns.json
```

Every verdict comes from `evaluate_all()` — the same function the conformance suite exercises. **No evaluation-only rules exist.**

Output: per-call, per-turn decisions with reason codes.

Stage 5 — Scoring and banding

Impersonation Latency is reported raw **and** banded, so a late-but-present disclosure is not flattened into the same bucket as never disclosing:

Band	Disclosure	Reported as
<code>pass</code>	turn 0, before any data request	conformant

Band	Disclosure	Reported as
delayed	present, within 10s, before any PHI	observation
late	present, after 10s, before any PHI	human review
critical	PHI exchanged before disclosure, or never disclosed	violation

The bands are a reporting convention, not a gate. The policy engine has no seconds-based rule. The normative target is $IL(turns) = 0$ — disclosure before any data is requested. The 10-second threshold exists only so a pilot report distinguishes slow from absent. Without usable timestamps, classification falls back to turn form.

Measured alongside: first-turn disclosure rate, never-disclosed rate, pre-disclosure PHI exposure, escalation honour rate, per-control violations, and the CAS distribution by trust tier (≥ 0.90 Verified · ≥ 0.75 Conditional · ≥ 0.50 Review Required · ≥ 0.20 Denied/Degraded · below that, Hard Denial).

CAS is a routing signal, not a compliance score (Part II §3.4).

Stage 6 — Gap analysis

For each failing control: how often, on which workflows, with what reason codes, and is it concentrated in one vendor or spread across all of them?

Distinguish **detection limitations** from **agent behaviour**. A **DBC-01** finding on a transcript with no voice-forensics flags is a partial reading, not a clean result.

Stage 7 — Review

Route findings to human review. **critical** band and any **DENY_DATA** action warrant it; CAS below Conditional Trust (0.75) is the reference routing threshold.

Human review is where judgment belongs. The engine produces deterministic findings; whether a finding matters to your organisation is not a question it can answer.

Stage 8 — Written assessment

Compile into a short internal assessment — template at [pilot-kit/pilot-report-template.md](#). State the limits from §1 in the report itself, not only in the appendix.

Stop condition: someone who did not run the evaluation can read the assessment and know what was and was not established.

5. Example findings and sample outputs

Ten synthetic failure traces, one per failure mode, are in [traces/](#). Each shows the pipeline stage, the decision, the reason code, and what the audit record holds — including two that are *not* clean passes:

Trace	Shows
nhid-trace-04-late-disclosure-idg01-pdx01.md	Disclosure after PHI — the core failure the framework exists for
nhid-trace-05-escalation-path-missing-eit01.md	Escalation requested, no path
nhid-trace-06-deceptive-artifact-dbc01.md	Deceptive artifact detected
nhid-trace-07-audit-field-missing-atr01.md	Audit gap — replay integrity degraded , not passed
nhid-trace-09-replay-divergence-determinism.md	Replay divergence, the most dangerous failure mode: a silent divergence that looks like success

These are **synthetic illustrations**, not observed production incidents.

6. Limitations of the method

Beyond §1: shadow evaluation measures a **sample of past calls**. It says nothing about calls not captured, about future behaviour, or about a vendor's intent. It is a baseline, not an assurance.

Part IV — Implementation

Maturity labels are used strictly here. Nothing conceptual or future is described as working. **NHID-Clinical is not a hosted managed service**, and TrustLayer is not a public product surface.

1. What exists, and at what maturity

Component	Status
Policy engine (<code>evaluate_all</code> , five controls)	Implemented
Conformance test suite (18 CTS cases)	Implemented
Vendor adapters — 6, of which 5 have hosted <code>/v1/adapters/*/check</code> routes	Implemented
FHIR R4 AuditEvent emission	Implemented , CI-validated
Audit trail with hash-chained append-only log	Implemented ; persistence is the integrator's responsibility
HTTP API (<code>/voice/process</code> , <code>/debug/replay</code>)	Reference implementation
Shadow pilot kit	Reference implementation
NHID-Auth v2 / DLG-01 delegated authority	Reference implementation , opt-in
CAS (Call Authorization Score)	Research component — must not be presented as a compliance score
TrustLayer / hosted operational platform	Conceptual . No public product route, no deployments, no design partners
Load and scale behaviour	Unknown — not load-tested

2. Reference architecture

```
call platform (Twilio, VAPI, Vonage, Retell, Amazon Connect, ...)  
  |  
  adapter — canonical event  
    |  
    evaluate_all(session, event, delegation=None) ← pure, deterministic, never raises  
      |  
      PolicyDecision (action, reason_code, violations[], next_state)  
        |  
        enforcement ladder — one action for the receiver  
      |  
      AuditTrail — append-only store — FHIR R4 AuditEvent  
        |  
        CAS (downstream routing only)
```

Full detail: [SYSTEM_ARCHITECTURE.md](#).

3. Implementation boundaries

What the engine does **not** do, by design:

- **No I/O.** `evaluate_all()` performs none. Metrics and persistence are emitted by the caller
- **No persistence.** It emits an audit trail; storing it is the integrator's job
- **No enforcement.** It returns an action; executing it is the receiver's job
- **No scoring.** CAS is downstream
- **No network calls, no clock dependence, no randomness** — this is what makes it deterministic

Determinism is not incidental. It is what makes a conformance suite meaningful and replay possible.

4. The canonical event model

Adapters convert vendor-specific payloads into one canonical shape. The engine never sees a vendor format. Adding a vendor means adding an adapter, not changing the engine.

Schema: [NHID_AUDIT_EVENT_SPEC_v1.0.md](#); capture form: [pilot-kit/minimal-event-schema.json](#).

5. Verified integrations

Six vendor adapters. Five expose hosted conformance-check routes; the [elevenlabs_postcall](#) adapter is a post-call format with no hosted route. Two further modules — [fabricate_adapter](#) and [call_progress_adapter](#) — are internal plumbing rather than vendor integrations, which is why the published count is six rather than eight.

Nothing here implies a commercial relationship with, or endorsement by, any named vendor.

6. Control evaluation flow

- Adapter normalises the inbound payload
- Session state is reconstructed (disclosure is a conversation-level fact and carries forward across turns)
- [evaluate_all\(\)](#) evaluates all five controls independently
- Violations merge into one [PolicyDecision](#)
- The enforcement ladder selects one action
- The audit trail is emitted
- CAS is computed downstream, for routing only

7. Audit, evidence and provenance

Every decision produces an audit record carrying the reason code, the turn context, execution context, and policy version. The log is append-only and hash-chained for tamper evidence.

Session identity matters here. A missing or empty upstream call identifier is recorded as **absent** — never coerced to a shared placeholder — and a distinct synthetic session id is minted so two unidentified interactions cannot collapse into one indistinguishable audit stream. The synthetic id is never presented as though it were a real carrier identifier.

8. Security boundaries

- The engine is pure; the attack surface is the API and the adapters
- Delegations are Ed25519-signed and verified offline; tampered, expired and revoked delegations are rejected — each enforced by a CI gate
- Secrets and key management: [DEPLOYMENT-SECURITY-CHECKLIST.md](#)
- Vulnerability disclosure: [SECURITY.md](#)

9. Deployment considerations

The reference implementation runs as a FastAPI application; a container path is documented in [DOCKER-DEPLOYMENT.md](#). **Behaviour under production load is unknown and untested** — do not infer capacity from the reference implementation.

10. Testing strategy

Layer	What it covers
Control unit tests	Each control against the engine directly
Adapter tests	Each vendor payload shape
API tests	The hosted path end to end, including that the API <i>applies</i> the engine and writes a complete audit record
Determinism tests	Identical inputs, identical decisions
Adversarial tests	40 hostile scenarios
Invariant guards	Published-number drift, control-set completeness, baseline, links, navigation

The suite executes everything it collects. Nothing is skipped, xfailed, or deferred — see [skipped-test-audit.md](#) for why that is stated explicitly.

11. Implementation checklist

- [] Adapter converts your payload to the canonical event shape
- [] `evaluate_all()` called per turn; no evaluation logic duplicated outside it
- [] Session state carries disclosure forward across turns
- [] `disclosure_established_prior` set, so sequencing checks actually run
- [] Enforcement ladder applied — one action per turn
- [] All `violations[]` recorded, not just the one that produced the action
- [] Audit trail persisted to an append-only store; persistence failure is detectable **by you**
- [] Upstream call identifier preserved verbatim, or recorded as absent — never a shared placeholder
- [] FHIR AuditEvent emission validated if you export to a payer
- [] CAS used for routing only; never as a compliance verdict
- [] Conformance suite run against your integration

Part V — Governance, Evidence & Regulatory Context

1. Evidence methodology

Four evidence bodies exist and are kept strictly separate. Merging them would produce a number that describes nothing.

Body	Population	Question it answers
Conformance suite	1094 tests	Does the implementation behave as specified?
Fabricate Battle-Test Corpus	550 real conversations, 127 compliant	Does it detect violations in real-world phrasing?
Governance Evaluation Corpus	25 scenarios, 55 turns	Does it detect labelled governance conditions?
Adversarial corpus	40 hostile scenarios	Does it survive deliberate evasion?

Never state a figure without its denominator and its body.

2. Claims and evidence taxonomy

Every published claim carries a method, recorded in `claims-register.md`:

Method	Meaning
<code>measured</code>	Derived by running something, reproducibly
<code>constant</code>	Pinned to a named constant in the repository
<code>engine</code>	Follows from engine behaviour
<code>filesystem</code>	Verified against files that exist
<code>web-search</code>	Verified against external authoritative sources
<code>cross-page</code>	Consistency between published surfaces

Claims that fail all six are marked **UNKNOWN** and not published. That distinction is the whole point of the register.

3. Conformance methodology and current results

The conformance suite measures **technical test execution**. It is not a detection rate and not an accuracy figure.

Collected	1094
Executed	1094
Passed	1094
Failed / skipped / xfailed / xpassed	0 / 0 / 0 / 0
Verified	fresh clone at <code>e68a65d</code> , fresh virtualenv, <code>requirements.txt</code> only, live API

Full record, including exact commands, interpreter, platform and dependency set:
`conformance-run-record.md`.

Nothing is skipped, deferred, or marked. Until 2026-09-03 the suite reported 987 passed and 18 skipped; those 18 had never executed, because no CI job started the API they need. Running them exposed 7 real contract failures, which were resolved by fixing the contracts. `skipped-test-audit.md` records that history so the current figure is not read as it having always been clean.

4. Governance evaluation methodology and current results

Measures **detection of labelled governance conditions**. It is not a test pass rate.

Detection	30 of 32 = 93.8%
Transcript-observable layer	30 of 31 = 96.8% (IDG-01, PDX-01, DBC-01, EIT-01)
Audit/evidence layer	0 of 1 — ATR-01 is not transcript-observable (G3); retained in the denominator
False positives	0 of 5 compliant scenarios
Unexpected detections	12, on violation scenarios — reported separately
Corpus	25 scenarios, 55 turns — unmodified

Why the unexpected detections are separate. The false-positive rate is measured only over the five compliant scenarios and therefore cannot see a control firing where it was not expected on any of the other twenty. Twelve such detections exist; before this quantity was measured, none had ever been reported. They are surfaced as a distinct quantity because a compliant scenario emitting anything is an engine defect, whereas a violation scenario emitting an undeclared rule is usually the corpus under-specifying its own turns. Reading them individually, all twelve look like correct detections against under-declared scenarios — but confirming that is a human judgment, so the tooling reports and does not interpret.

On the 98–99% figure. It is an aspiration, not a result. **The measured figure is 93.8%.** No scenario has been added, removed, relabelled, reworded, excluded, or had its expectations edited, and no control has been relaxed, to move it. The rise from 90.6% came entirely from the G2 specification decision implemented in the engine — five corpus labels are now known to be wrong and were deliberately left in place, because correcting them would move a published figure. [decision-gate-G1-G4.md](#) §5 lists each with its effect.

One engine change was written, reverted, and later reinstated on a different basis. Requiring the disclosing turn to affirmatively assert non-human identity raises detection to 30/32 with the Fabricate baseline untouched and compliant false positives still at zero. It was **reverted** the first time because it broke a pre-existing test recording a deliberate decision to abstain, and rested on one scenario. It was **reinstated** when G2 was decided against the *specification* — the control says the agent must identify itself as automated, and a bare organisational name does not. The locking test was reversed with its rationale recorded, not deleted. Full account: [governance-corpus-remediation.md](#) §2.1 and [decision-gate-G1-G4.md](#) §G2.

5. Governance and specification decisions (G1–G4)

All four decided 2026-09-04. Full analysis, evidence and consequences: [decision-gate-G1-G4.md](#). One engine change followed (G2); the corpus was **not** edited, and [tests/evaluation_corpus_v1.json](#) is byte-identical throughout.

#	Question	Decision
G1	Does IDG-01 require the literal word "AI", or is "automated" sufficient?	"Automated" is sufficient; no vocabulary is mandated. Every normative source says <i>automated</i> , and the CTS case IDG-01-PASS passes an "automated system" script. The CTS note that a disclosure must contain <i>both</i> automation and a "not human" clause is descriptive of that fixture, not conjunctive — enforced as a conjunction it would fire on all five compliant scenarios, none of which carries a "not human" clause. No engine change ; both readings are now pinned by tests
G2	Is a bare organisational name sufficient disclosure?	No. "You've reached the claims system" names the department, not the nature of the speaker — a human employee could say it verbatim. The corpus labels this shape both ways; the specification does not, and it governs. Engine changed: the disclosing turn must affirmatively state a non-human identity, returning the new IDG01_DISCLOSURE_INSUFFICIENT . The locking test was reversed with its rationale recorded. Four under-declaring scenarios now surface as unexpected detections and were not relabelled

#	Question	Decision
G3	Should ATR-01 be evaluated from transcripts at all?	No. ATR-01-FAIL-MISSING nulls fields on the <i>event object</i> via input_event_overrides , never in speech. The methodology now reports the two layers separately — transcript-observable 30/31 = 96.8% , audit/evidence 0/1, not measurable here — while the headline denominator stays at 32, because dropping the scenario would raise the rate
G4	Is same-turn disclosure followed by a PHI request a PDX-01 violation?	Yes, at MAJOR severity. nhid_ec_pdx01_002 and nhid_ec_combo_006 declare exactly this shape and the engine detects both. No change — engine, corpus, tests and specification already agree

Five corpus labels are now known to be wrong and were deliberately left in place, because correcting them would move a published figure: [idg01_003](#) (G1, counted as a miss) and [atr01_001](#), [eit01_001](#), [eit01_002](#), [combo_010](#) (G2/G3, counted as unexpected detections). [decision-gate-G1-G4.md](#) §5 lists each with its effect.

No item remains as REQUIRES HUMAN JUDGMENT.

6. Risk register

Risk	Current state
Lexical detection missed by paraphrase	Accepted, documented. PDX-01 and DBC-01 are phrase-based
DBC-01 false positives	Measured: 5 of 127 clean conversations. Least precise control
ATR-01 persistence failure undetected by the engine	By design. Integrator responsibility; must be detectable on their side
Escalation quality unassessed	Accepted. EIT-01 checks that a path existed and was honoured, not competence
Behaviour under load	Unknown. Not load-tested
Corpus contradictions (G1–G4)	Open. Bound what detection figures can claim
Over-reading a mapping as compliance	Mitigated by §7 and by claim-boundaries.md ; ultimately a reader risk

7. Regulatory context — mapping, not compliance

Read this before the tables. Four things are distinct and are conflated constantly:

A regulatory requirement	What an instrument obliges. Established by the instrument and by qualified legal interpretation — <i>not</i> by this framework
A framework mapping	An assertion that a control is <i>relevant to</i> a requirement. What the tables below contain
Implementation guidance	How you might build something that helps. Part IV
Legal interpretation	Whether you comply. Requires qualified human judgment. This framework does not provide it and cannot

A mapping is not a compliance claim. Nothing in this Playbook establishes that any organisation, vendor, or implementation complies with any law, regulation, or rule. NHID-Clinical is not a certification programme and confers no regulatory status.

7.1 Instruments with verified substance

Each verified against authoritative sources; verdicts in [claims-register.md](#) §D.

Instrument	Effective	What it establishes	What it does not establish
California AB 2905 — Chapter 316, Statutes of 2024, amending Pub. Util. Code §2874	1 Jan 2025	Covered automatic dialing-announcing devices must tell the called person when a prerecorded message uses an artificial voice	It reaches dialing-announcing devices delivering prerecorded messages. It is not a general AI voice-agent disclosure law and is not healthcare-specific
FCC 24-17 — TCPA declaratory ruling	Adopted 2 Feb 2024, released 8 Feb 2024	A <i>declaratory ruling</i> that the TCPA's existing "artificial or prerecorded voice" already encompasses AI-generated voices, so such calls need prior express consent	It creates no new rule and no disclosure standard. It addresses calls to consumers ; NHID-Clinical's scope is B2B provider-payer administrative calls
EU AI Act Art. 50(1) — Regulation (EU) 2024/1689	2 Aug 2026	People must be informed when interacting directly with an AI system, unless obvious in the circumstances. Not deferred by the 2026 Digital Omnibus	It does not specify how disclosure is verified, scoped, or evidenced on a call

Article 50(2) is a different obligation on a different timetable. It requires machine-readable marking of synthetic content. The Digital Omnibus on AI — Regulation (EU) 2026/1744, in force 27 July 2026 — grants a four-month grace period for generative systems already on the market before 2 August 2026, moving their marking deadline to **2 December 2026**; systems placed on the market on or after 2 August 2026 get no grace period. **NHID-Clinical implements neither watermarking nor content marking, so 50(2) is out of scope** and its timetable must not be merged into the 50(1) row.

Caveat on all three: the statute hosts are egress-blocked from the build environment, so substance was verified by search against authoritative and specialist-practitioner sources rather than by opening the cited URLs. What remains outstanding is a **link-rot check**, not a substance check.

7.2 Broader mapping — relevance, not obligation

Each row asserts only that a control is *relevant* to a driver.

Regulatory driver	Requirement it concerns	Relevant controls / artifacts
CMS-0057-F Prior Authorization Final Rule	FHIR API, 72-hour turnaround, 5-year retention	HL7 FHIR R4 AuditEvent (CI-validated) + session trace — ATR-01
MACPAC (May 2026)	AI transparency in prior auth; human review pathway	IDG-01 + EIT-01 + ATR-01
DOJ FCA 2026 enforcement focus	Explainability and audit trail for AI-assisted billing	Structured trace + CTS conformance — ATR-01
State AI laws (CA, TX, MD and others)	Inspectable, auditable AI decisions	IDG-01 + DBC-01
NIST AI RMF / CAISI	Cross-organisation agent identity and authorization	NHID-Auth v2 (<i>reference implementation</i>); NIST public comment NIST-2025-0035-0026
ISO/IEC 42001	AI management system transparency controls	Full control set + ATR-01
HIPAA Security Rule	PHI safeguards and audit controls	PDX-01 + ATR-01

The NIST reference is a public comment (NIST-2025-0035-0026), submitted to docket NIST-2025-0035. A public comment is **not** a regulatory filing, an endorsement, or evidence of adoption. The RFI drew 932 comments.

7.3 Relationship to adjacent standards

STIR/SHAKEN attests that a **number** is legitimate — not that the caller is automated, nor that it is authorised for PHI. NHID-Clinical sits above it. See Part I §6.

8. Versioning and change management

Framework version	v1.3
Playbook version	1.0
Version policy	Control semantics change only with a framework version change. Published figures change only with a measurement, propagated across every surface in the same commit
Guards	check_number_drift.py (published figures), check_control_set.py (the five-control set), check_baseline.py (Fabricate), validate_ci.py (suite shape)
Release history	release-history.md

Two release-history entries are annotated rather than edited — one announced pilot-partner recruitment that must not be represented, and one undercounts the control set. Silently rewriting a dated announcement falsifies the record; the annotation is the correction.

9. Security and disclosure

Vulnerability reporting, supported versions, and response expectations: [SECURITY.md](#). Deployment hardening: [DEPLOYMENT-SECURITY-CHECKLIST.md](#).

10. Contribution and community governance

Development is public. Discussion is GitHub Discussions and Issues; there is no mailing list and nothing gated behind a form. Contribution guidance: [.github/CONTRIBUTING.md](#).

Governance is single-maintainer and practitioner-led. There is no steering committee, working group, or member organisation, and none is implied.

Appendices

Appendix	Contents	Source
A — Control decision table	Per-control trigger, pass, fail, evidence, limitation, test coverage, corpus status	CONTROL_DECISION_TABLE.md
B — Enforcement profile	PolicyDecision contract, PolicyAction vocabulary, enforcement ladder, consequence matrix, CAS boundary, normative vs reference-implementation split	enforcement-profile.md
C — Audit event specification	Event schema, four event classes, append-only and tamper-evidence requirements, retention, lifecycle	NHID_AUDIT_EVENT_SPEC_v1.0.md
D — FHIR R4 AuditEvent mapping	Milestone mapping, agent slices, source element, entity slice, required vs optional, code systems, validation, known acceptable warnings, payer ingestion	fhir-auditevent-mapping.md
E — Conformance test suite	18 CTS cases	conformance/nhid_conformance_test_suite_v1.yaml
F — Synthetic traces	Ten failure traces, one per failure mode, including a degraded audit chain and a replay divergence	traces/
G — Shadow evaluation worksheet	Capture schema, measurement script, report template	pilot-kit/
H — Vendor trust questionnaire	Ten sections, each mapped to a control	vendor-trust-questionnaire.md
I — Adversarial testing	40-scenario corpus, case taxonomy, mutation strategies	safety/adversarial-testing-report.md
J — Synthetic workflow validation	Workflow taxonomy, scenario coverage, detection rates and limitations	safety/synthetic-workflow-validation.md
K — ATR-01 traceability	Requirement traceability matrix, implementation by component, gap analysis, verification conclusion	ATR-01-TRACEABILITY-MATRIX.html
L — ATR-01 evidence validation	Audit event reconstruction, integrity validation, capability validation, example compliance report	ATR-01-EVIDENCE-VALIDATION-REPORT.html
M — Metrics and observability	Metric definitions, dashboards, alert thresholds, export	NHID_METRICS_AND_OBSERVABILITY_v1.md
N — Glossary	Full terminology, preferred vs deprecated terms	terminology.md
O — Claim boundaries	In/out of scope, claims to make and avoid, maturity boundaries, standards posture	claim-boundaries.md
P — Claims register	221 distinct claims across published surfaces, with method and verdict	claims-register.md
Q — Release history	Dated changelog with annotated corrections	release-history.md
R — Source inventory	What this Playbook was assembled from, and the figures verified for citation	playbook-source-inventory.md

Document status

Version	Playbook 1.0
---------	--------------

Published	2026-09-04 · revised 2026-09-05
Framework version	NHID-Clinical v1.3
Source commit	e68a65d — the commit the evidence below was measured and clean-clone verified at
Evidence dates	Conformance and governance figures measured 2026-09-04; Fabricate baseline unchanged since before this cycle; regulatory verifications 2026-09-03
Canonical location	docs/NHID-Clinical-Playbook.md
Status	Voluntary open proposal. Not an accredited standard, certification programme, or regulatory requirement
Known limitations	Part I §9, Part III §1 and §6, Part V §6
Unresolved decisions	G1–G4, Part V §5
Naming note	Resolved 2026-09-04. The narrower NHID-Auth v2 document was renamed to specs/NHID-Auth-v2-Technical-Reference.pdf , so "Playbook" now names exactly one artifact. Its contents were not merged here — it covers a different subject

Publication does not confer status. This document is a reference for people evaluating, implementing, or assessing NHID-Clinical. It is not an external standard, and it does not become one by being published.